

UNIVERSITY OF TARTU
FACULTY OF MATHEMATICS AND COMPUTER SCIENCE
Institute of Computer Science

LAURIS KRUUSAMÄE

Optimizing Contextual Sensor Data Collection for Mobile Users

Bachelor Thesis (6 EAP)

Supervisor: Satish Narayana Srirama, PhD
Co-supervisor: Huber Flores, MSc

Author:..... "...." Month 2013

Supervisor:..... "...." Month 2013

Professor:..... "...." Month 2013

TARTU, 2013

Abstract

Nowadays, mobile applications are becoming more context aware due to technological achievements which enable the applications to anticipate users' intentions. This is achieved through using the device's own micromechanical artifacts that can be used to perceive the environment. However, this is constrained to the hardware limitations of devices.

A proposed solution for this has been made in the thesis "Context Sensor Data on Demand for Mobile Users Supported by XMPP" by Kaarel Hanson. The solution is to use XMPP for transporting sensor data from Arduino microcontroller to the cloud. Arduino provides low-cost hardware, while the cloud offers the reliable and high-availability means for storing and processing sensor data.

This solution shows that running on a 9V battery the microcontroller lasts for 101 minutes when using an Ethernet module for communications, and 161,5 minutes with a WiFi module. These results are not good enough for remote data collection with limited access to the microcontroller.

This thesis proposes an optimisation for the system so that instead of reading and sending sensor data every 10 seconds, the cloud server would notify the controller when to start sending data and when to stop. This means implementing an algorithm for detecting similar sensor data readings and notifying the microcontroller of needed operations. With similar readings, the microcontroller could be put to an idle state for limiting power consumption, which would prolong battery life.

The aim is to optimise the sensor reading process enough to prolong the Arduino microcontroller's battery life on a 9V battery.

Contents

List of Figures	v
1 Introduction	1
1.1 Introduction	1
1.1.1 Motivation	1
1.1.2 Contributions	1
1.1.3 Outline	2
2 State of the Art	3
2.1 Arduino	3
2.1.1 Arduino Mega ADK	4
2.1.2 Wireless SD Shield	5
2.1.3 RN-XV WiFly Module	6
2.1.4 TinkerKit	7
2.2 Fuzzy Logic	7
2.2.1 Fuzzy Set	8
2.2.2 Fuzzy Set Operations	9
2.2.3 Fuzzy Control Systems	10
2.2.3.1 Fuzzy control process	11
2.3 Simple Linear Regression	12
3 Problem Statement	13
3.1 Current Solution	13
3.1.1 Arduino	13
3.1.2 XMPP Communication	14
3.1.3 Data Collection Server	15

CONTENTS

3.2 Problems	16
3.2.1 Hardware	16
3.2.2 Software	17
3.2.3 Power Consumption	17
4 Problem Solution	18
4.1 Power Consumption	18
4.1.1 Sleep	19
4.1.1.1 Watchdog Timer	19
4.1.1.2 Arduino Mega ADK	19
4.1.1.3 WiFly module	20
4.1.2 Communication	20
4.2 Server-side Client	21
4.2.1 Simple Linear Regression	23
4.2.2 Fuzzy Logic Engine	24
4.3 Results	28
5 Conclusions	29
5.1 Conclusions	29
5.2 Summary of Contributions	29
6 Related Work	30
7 Future Research Directions	32
8 Sisukokkuvõte	33
Bibliography	34

List of Figures

2.1	Arduino Mega ADK	4
2.2	Wireless SD Shield	5
2.3	RN-XV Wifly Module	6
2.4	Complete Arduino Sensor Module	8
2.5	Trapezoidal and Triangular Membership Functions	9
3.1	XMPP Session Lifecycle (1, p. 40)	15
3.2	Server-side Data Model	16
4.1	Power Consumption at 30 Second Interval	22
4.2	Updated Server-side Data Model	23
4.3	Regression Confidence Fuzzy Sets	25
4.4	Sensor Value Fuzzy Sets	26
4.5	Idle Time Fuzzy Sets	27

1

Introduction

1.1 Introduction

1. Write this part Briefly summarize the question (you will be stating the question in detail later), and perhaps give an overview of your main results. (it is not just a description of the contents of each section)

1.1.1 Motivation

The developed prototype (1) works well when the general idea and goal is considered. However, the lack of flexibility of data transmission intervals and high power needs constrain the implementation usages. In order to take full advantage of the prototype, power usage should be reduced and more flexibility added to the communication between the server-side client and sensor modules.

1.1.2 Contributions

A more flexible data collection solution is developed based on HTTP. The changes made enable more energy-efficient data collection by predicting sensor data when possible. A fuzzy control system was developed to decide if sensor data is predictable and for what time period. Moreover, the prototype's Arduino-based sensor module has been improved to support sleep mode while no sensor data is sent to the server-side client.

1.1 Introduction

Power consumption before and after the changes was tested to indicate changes' effect on battery life. **2. mention peaktech power supply here?**

1.1.3 Outline

Chapter 2: introduces the Arduino framework and modules used in the prototype. Finally, a description of fuzzy control systems and simple linear regression is given.

Chapter 3: describes the existing prototype and gives an overview of its implementation details. Later, an overview of identified areas of improvement is given.

Chapter 4: explains the improvements made to the prototype. Firstly, describes changes made to the Arduino sensor module. Secondly, changes in the communication between the sensor module and the server-side client are discussed. Lastly, an overview of the server-side fuzzy control system and data prediction techniques is given. The chapter is ended by a power consumption comparison between the initial and improved prototype.

3. add more chapters here when they're ready

2

State of the Art

The state of the art used in the thesis highlighted the advances in the cloud computing domain and the mobile domain...

Description of Jabber and XMPP protocol and its usage.

2.1 Arduino

Arduino (2) is an open-source electronics prototyping platform based on a simple microcontroller board and a development environment for writing software for the board. It is intended for anyone interested in creating interactive solutions. Arduino can take inputs from a variety of sensors and control various actuators and lights. The microcontroller is programmed using the Arduino programming language (based on Wiring) and the Arduino development environment (based on Processing). Arduino IDE enables to choose between different board models, microcontroller programmers and communication ports.

4. add an image of the IDE here... Programs (called sketches) are written and uploaded to the board using the Arduino IDE. Each sketch must have two functions- *setup()* and *loop()*. *setup()* is the first function called after Arduino is started or rebooted. It is called once and afterwards the function *loop()* is called consecutively until the board is stopped, restarted or crashes. When a crash occurs, the program is restarted, which means calling *setup()* again.

Since programs are written in C/C++, there are a lot of libraries available for use.

2.1.1 Arduino Mega ADK

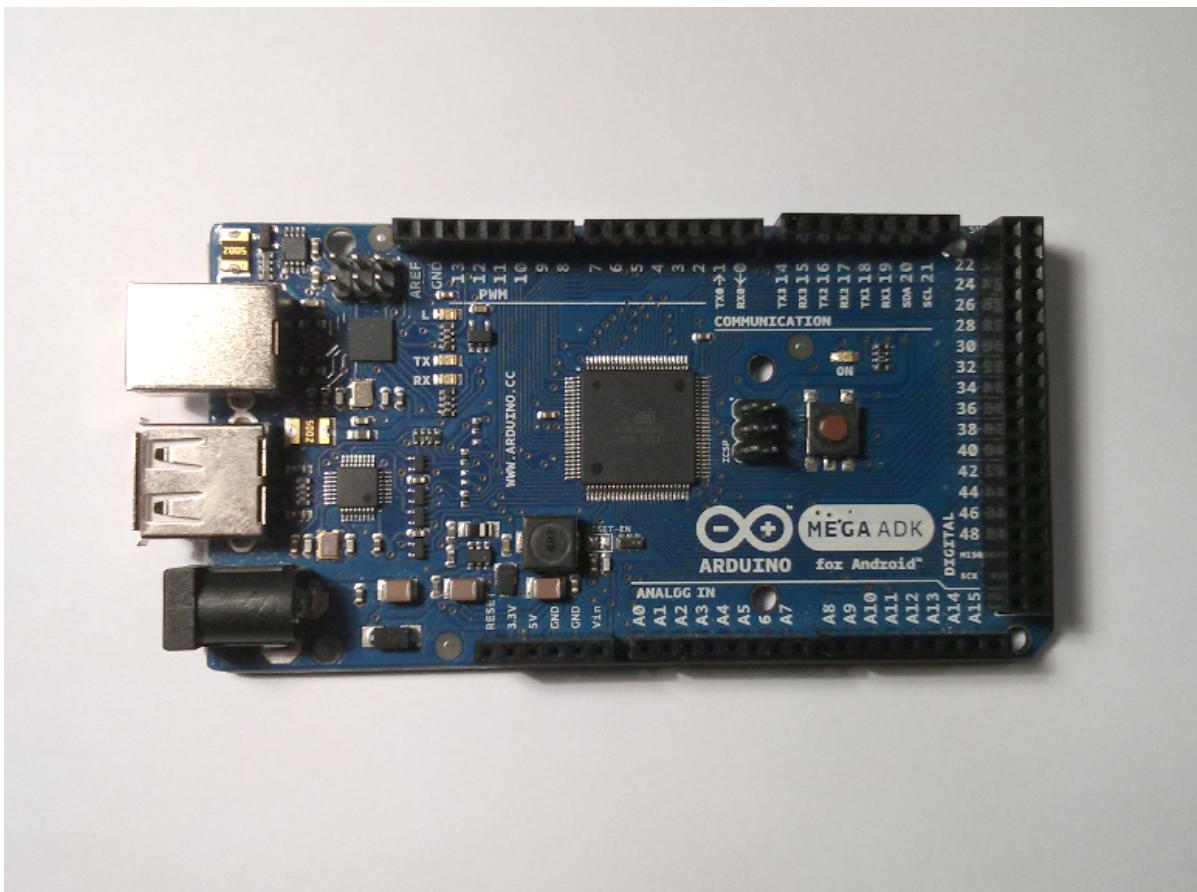


Figure 2.1: Arduino Mega ADK

Arduino Mega ADK (3) is one of the most capable boards available. The Arduino ADK is a microcontroller board based on the ATmega2560. Similar to the Mega 2560 and Uno, it features an ATmega8U2 programmed as a USB-to-serial converter. It has a USB host interface to connect with Android based phones, 54 digital input/output pins, 16 analog inputs, 4 UARTs (hardware serial ports), a 16 MHz crystal oscillator, USB B, micro B connections and a 2.1mm center-positive power jack. The ADK is designed to be compatible with most shields designed for the Uno, Diecimila or Duemilanove

The ADK has 256 KB of flash memory for storing code (of which 8 KB is used for the bootloader), 8 KB of SRAM and 4 KB of EEPROM (which can be

2.1 Arduino

read and written with the EEPROM library).

The Arduino ADK can be powered via the USB connection or with an external power supply. The power source is selected automatically. External (non-USB) power can come either from an AC-to-DC adapter or battery. The board can operate on an external supply of 5.5 to 16 volts. The recommended range is 7 to 12 volts. The board can be seen in Figure 2.1

5. format the text a bit, remove describe what is EEPROM, add an image

2.1.2 Wireless SD Shield



Figure 2.2: Wireless SD Shield

The Wireless SD shield (4) allows an Arduino board to communicate using a wireless module. The module can communicate up to 100 feet indoors or up to

2.1 Arduino

300 feet outdoors. The module can be seen in Figure 2.2

The shield has an on-board switch which allows to select between USB and Micro modes. In USB mode, the shield bypasses Arduino board's microcontroller and communicates directly to the USB-to-serial converter. In Micro mode, data sent from the microcontroller will be transmitted to the computer via USB as well as being sent wirelessly by the wireless module. The microcontroller will not be programmable via USB in Micro mode.

2.1.3 RN-XV WiFly Module




Figure 2.3: RN-XV Wifly Module

The RN-XV module (5) by Roving Networks is a certified Wi-Fi solution. It is based upon Roving Networks' robust RN-171 Wi-Fi module and has 802.11 b/g

2.2 Fuzzy Logic

radio, 32 bit processor, TCP/IP stack, real-time clock, power management unit. It has a pre-loaded Roving firmware to enable simple configuration and usage. The module can be seen in Figure 2.3

2.1.4 TinkerKit

TinkerKit (6) is a tool used to build interactive products using Arduino boards. It consists of modules (sensors, actuators) and a sensor shield. The tool greatly simplifies product assembly, because instead of building circuits out of low level components, all the modules can be attached to the TinkerKit sensor shield with a snapping cable.

The modules are divided into sensors and actuators. Sensors can measure their surrounding environment (provide input) and actuators perform actions (provide output). The sensors used are thermistor, LDR and Hall. The fully assembled module is shown in Figure 2.4

Thermistor sensor (7) is a resistor whose resistance changes significantly with temperature. The module's output ranges from 0V to 5V

Hall sensor (8) creates a voltage related to the magnetic field near the sensor. It can detect the presence of a nearby magnet or magnetic field induced in a coil or wire. The output ranges from 0V (no presence) to 5V (presence detected).

Light Dependent Resistor or LDR (9) is a variable resistor whose resistance is decreased when light falls on the sensor. The module's output is 5V when the module receives no light and 0V when bright light falls on it.

2.2 Fuzzy Logic

6. add citations here...

Fuzzy logic is a form of logical reasoning which is based on approximate instead of exact reasoning. When traditional reasoning has binary values like true (1) and false (0), fuzzy logic has values ranging from 0 to 1. These values represent a partial truth or in other words the degree of truth. This is similar to probability which also describes partial truth. The difference between probability and fuzzy

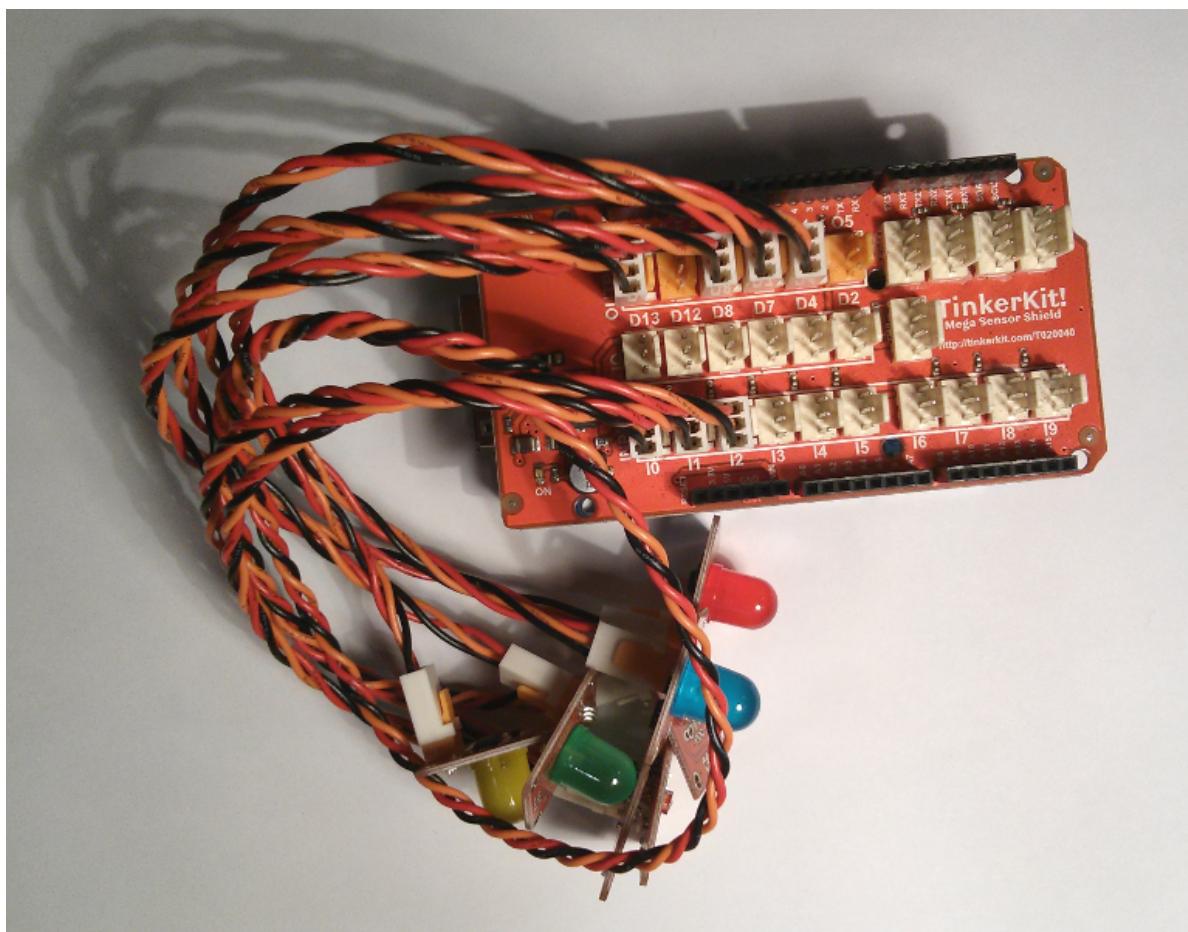


Figure 2.4: Complete Arduino Sensor Module

logic is that whereas fuzzy logic describes how true something is, probability describes how likely it is for something to be true.

2.2.1 Fuzzy Set

Fuzzy sets are the main building blocks of fuzzy logic. In traditional set theory where an element can either belong or not belong to a set. With fuzzy sets, an element has a degree of membership to a set.

Fuzzy sets are represented by membership functions, which measure the degree of membership (DOM) a given value has to the set. Membership functions are usually depicted graphically and have names based on the shape of their

2.2 Fuzzy Logic

graphical representation. Membership functions are depicted as 2-dimensional graphs, where the x-axis represents the values and y-axis the degrees of membership (ranges from 0 to 1). The degree of membership for a given value can be found by projecting vertically to the upper boundary of the membership function.

The most used membership functions are triangular and trapezoidal functions. Sample sets are shown in Figure 2.5. Based on these graphs the DOM can easily be seen. Furthermore, from this graph can be seen the difference between fuzzy sets and traditional sets. Let us assume that the values represent age, trapezoidal membership function describes young people and triangular describes middle-aged people. Now when we have an age value that is between points X_2 and X_3 , this value belongs to both the young and middle-aged sets. The difference is in the value of DOM to either set. When the value is nearer to point X_3 than to X_2 the statement that it belongs to the middle-aged set has a higher truth value.




Figure 2.5: Trapezoidal and Triangular Membership Functions

2.2.2 Fuzzy Set Operations

Fuzzy set operations (10) are a generalized version of standard set operations. The most widely used generalization is referred to as the standard fuzzy set

operations. This generalization includes three main operations - complement, union and intersection. Lets assume we have two fuzzy sets with membership functions A and B , then these operations can be defined as:

- Complement:** $\neg A(x) = 1 - A(x)$
- Intersection:** $(A \cap B)(x) = \min\{A(x), B(x)\}$
- Union:** $(A \cup B)(x) = \max\{A(x), B(x)\}$

2.2.3 Fuzzy Control Systems

Fuzzy control systems (11) take a heuristic modeling approach to describing the controllable domain. The heuristic model is based on a set of rules in the form: **if** *< condition >* **then** *< action >*.

With the fuzzy logic part, the control system has a set of input variables and output variables, which are mapped into fuzzy sets. These variables and fuzzy sets have linguistic variables called labels.

These fuzzy sets and rules combined gives us a control system which consists of a set of rules where conditions and actions are defined as fuzzy values. For example if we have an input variable *temperature* with a fuzzy set labelled *cold* and an output variable *heat* with a fuzzy set labelled *on*, we could write a following rule: IF *temperature* = *cold* THEN *heat* = *on*. Based on this rule, the system can decide that if the temperature is cold then the heat should be turned on.

In an actual system there are multiple rules defined. As rules are activated only if their condition holds, the process of adding rules and variables has little computational overhead making the rulebase easily extendable.

Fuzzy control systems are mostly based on experience and experiments rather than actual mathematical modelling. Furthermore, as input and output variables are mapped into fuzzy sets, which have linguistic labels, the systems control rules are easily understandable. This makes fuzzy control systems usable in cases where the actual users have little knowledge of the implementation details. Instead, they can focus on the control process by defining understandable variables, their fuzzy sets and rules based on these sets.

2.2.3.1 Fuzzy control process

The decision making process has 4 main steps - fuzzification, rule evaluation, interference and defuzzification. These steps assume that the rules, input and output variables have already been defined. of the model

The first step in the process is fuzzification. In this step, the system's input and/or output variables are mapped into fuzzy sets. Each input variable can have one or more associated fuzzy sets. What this process does is decompose the crisp input values into fuzzy values. These fuzzy values are expressed by linguistic terms (labels) which can then be used in the rule evaluation process.

The next step is to use the fuzzified values in rule evaluation. A set of pre-defined rules is evaluated based on the fuzzy values created during the fuzzification process. Rules are activated when their condition holds, otherwise they are skipped. After all rules have been considered, the activated ones are gathered together. The results or actions of these rules are the basis of the next step called defuzzification.

Defuzzification is the final step in the fuzzy control process. This part is responsible for translating the fuzzy values (linguistic terms) into a crisp quantifiable result. Usually in the previous step, rule outputs will be linguistic terms (fuzzy sets) and corresponding degrees of membership for a defined output variable. In the defuzzification part, the different fuzzy sets and their DOMS for the output variable must be added together somehow to produce a final output.

There are several defuzzification methods. The simplest of which would be to use the fuzzy set with the highest DOM as the final result. However, center of gravity is the most commonly used method. First, the different degrees of membership calculated from rules must be combined for each output variable's fuzzy sets. Then, the membership functions are depicted of a graph and for each fuzzy set, the top of the membership function is cut off in a straight horizontal line at the height of the DOM value. When this is done for all fuzzy sets of the output variable, the center of gravity for the remaining shape is calculated. The center point's projection on the value axis (usually the x-axis) is the crisp result.

2.3 Simple Linear Regression

2.3 Simple Linear Regression

7. Add simple linear regression here, too?

3

Problem Statement

In this chapter, the prototype developed in Context Sensor Data on Demand for Mobile Users Supported by XMPP (1) is described. The prototype was successful, but had some areas of improvement. Secondly, an overview of these improvements is given.

8. Add a small introduction of who wrote the previous thesis etc

3.1 Current Solution

The current solution (1) has three main components:

1. Arduino sensor module
2. OpenFire XMPP server
3. Data collection server in the cloud (referred to as the server from here on)

There were two separate configurations described in Context Sensor Data on Demand for Mobile Users Supported by XMPP - one using Wi-Fi and the other Ethernet for network communication. Only Wi-Fi configuration is considered in this thesis due to the fact that the availability of an Ethernet connection (a cable) usually means that there is a power outlet nearby.

3.1.1 Arduino

The first component is the Arduino sensor module. The hardware configuration is based on the Arduino Mega ADK board. Wireless Shield with RN-XV WiFly

3.1 Current Solution

module is mounted on top of the board. Wireless Shield and Arduino board communicate over UART (hardware serial). TinkerKit Mega Sensor Shield is mounted on top of the Wireless Shield with 7 modules attached to it: 4 LED indicator lights, Hall, thermistor and LDR sensors. The external power source used is a 9V battery.

On the software side, the implementation mainly relies on a XMPP library and WiFi module library called WiFiHQ. The WiFiHQ library is responsible for creating a TCP connection and communication over the network, the XMPP library handles all the XMPP implementation details.

The sketch itself is fairly straightforward - in *setup()* a wireless network is joined, a TCP connection established and lastly a XMPP session is initialized. In *loop()* all connections are checked and reestablished if needed. Then the last transmit time variable is compared to the current time and if the report step amount (currently 10 seconds) has passed, data is collected from the sensors, formatted to appropriate JSON string and sent to the server-side client. A sample data JSON:

```
{"location": 1, "data": [{"type": 1, "value": 5.00}, {"type": 2, "value": 500.00}, {"type": 3, "value": 312.00}]} 
```

3.1.2 XMPP Communication

Both the data collection server and Arduino module are XMPP clients. The XMPP OpenFire server runs in the cloud and provides XMPP communications to both clients. Clients connect to the same chat where the server listens for messages from the sensor module. When a message is received by the server, sensor data is parsed from it and saved in a database.

XMPP session lifecycle can be seen in Figure 3.1. With the current implementation, steps 1 - 5 are done once when establishing the connection or when the connection drops. Data transmission step is done every 10 seconds and the last 2 steps are done when the connection is closed.

3.1 Current Solution

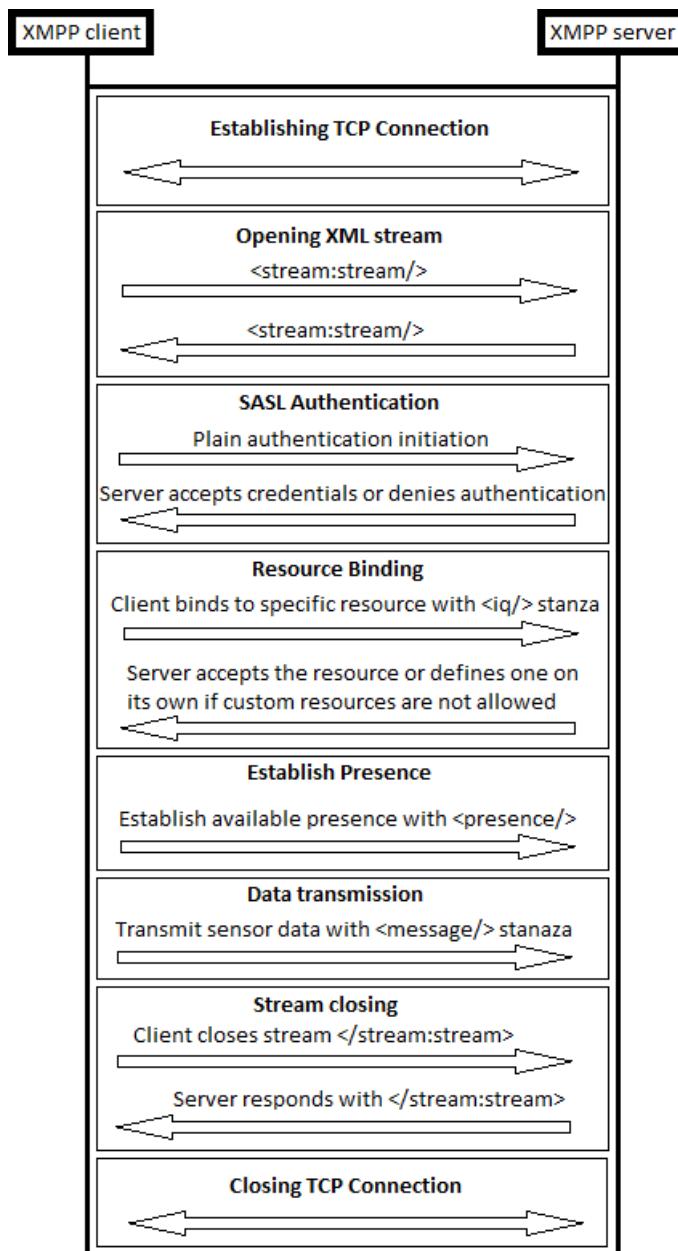


Figure 3.1: XMPP Session Lifecycle (1, p. 40)

3.1.3 Data Collection Server

The data collection server is responsible for gathering sensor data and saving it in a database. Its implementation is written in Java and uses Smack XMPP

3.2 Problems

API to communicate over XMPP. Data is stored in a H2 database, because of its simplicity and suitability for prototyping.

The database has three main tables - sensors, data and locations. Locations table has different sensor module locations, sensors table has different types of sensors and data has data gathered from different locations and sensors. The data model can be seen in Figure 3.2.

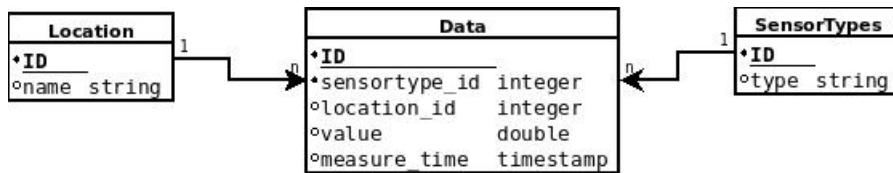


Figure 3.2: Server-side Data Model

3.2 Problems

Previously mentioned implementation has a few problems which will be discussed in this section. The two main points are power consumption and data collection flexibility.

According to the tests run in the previous thesis, the battery is able to power the module for 161,5 minutes using Wi-Fi. (1, p. 50-51) This, however, is not sufficient to enable actual data collection from a remote location. The problem can be addressed by using a larger battery, but the actual power consumption should be optimized, too. Additionally, the 10 second data transmission interval is hard coded into the Arduino sensor module, which does not provide enough flexibility. The following chapter described changes to the hardware, software and communication methods which improve on these error points.

3.2.1 Hardware

The hardware configuration has two main problems. Firstly, communication over the wireless module can be unstable at times, because parts of the messages might be missing or scrambled when read from the UART (1, p. 47). Since XMPP

3.2 Problems

session initialization is quite verbose, the possibility of receiving scrambled or incomplete messages effects the process.

Secondly, all parts of the hardware configuration are run in full power mode. As there is a 10 second gap between data transmission and sensor reads, there is a time period when most components are idling. This, in turn, means that some parts of hardware could use less power during these intervals and therefore reduce overall power consumption.

3.2.2 Software

With the existing implementation *loop()* is called consecutively and time differences are checked to determine when 10 seconds has passed using the internal *millis()* function, which gives the current Unix time. Since we have a 10 second interval, we do not need to waste cycles on time difference checks we now will fail for the next 10 seconds. A way to delay program execution could be used to stop the execution until we know the defined amount of time has passed.

In addition, as messages from the Wireless Shield might not be complete and XMPP session initialization is a verbose process, the current XMPP implementation hangs when scrambled messages are received during session start up. The library checks from complete XML tags, but when the attributes are incomplete, connection will not be successful and the session initialization code hangs.

3.2.3 Power Consumption

Power consumption is the main problem this thesis focuses on. As a 9V battery has an average capacity of 400 *mA*h at 100 *mA* current, the resource is quite limited. With the existing configuration, the average current drawn is 108 *mA*. Battery lifetime tests in the previous thesis showed that the module runs for 161,5 minutes on a 9V battery (1, p. 50). However, 161,5 minutes is not enough to gather contextual data for the proposed data collection system. The batteries need to be changed too often for it to be a viable solution.

9. how should i

reference all the data collected?

4

Problem Solution

In this chapter improvements to the existing implementation are discussed. First, changes in the Arduino sensor module are described. Secondly, an overview of changes in the communication between the module and data collection server (the server from here on). Lastly, an overview of the fuzzy control and data prediction systems that were implemented to reduce communication overhead and power consumption.

The main goal is to reduce the average power consumption measured by [10.](#) [reference it somehow](#) the PeakTech power supply.

[11. introduction](#)

4.1 Power Consumption

The first problem addressed is power consumption. There are two main ideas behind reducing it - put the sensor module to sleep mode when not transmitting data and reduce the need for data transmission. For this, the server-side client was improved to predict sensor data when possible and notify the sensor module of the next data transmission time. This enables the module to enter sleep mode for the given time period and thus reduce power consumption.

4.1 Power Consumption

4.1.1 Sleep

Since the module consists of 3 components - Arduino Mega ADK, Wireless SD Shield with WiFly module and TinkerKit Sensor Shield. Both the Mega ADK and WiFly module support sleep modes.

4.1.1.1 Watchdog Timer

A watchdog timer (12) is an electronic timer used to recover from computer malfunctions. They are found in automated systems where human interference is not possible and therefore the system must be able to recover from malfunctions on its own. A watchdog timer essentially performs a timing function producing a delayed response to an input trigger. The most common implementation has a digital counter that counts from a specified value down to a terminal value. Usually the initial value is programmable. When the counter reaches the terminal value, the timer timeouts and triggers a timeout signal. Usually this means restarting the program from the start. A program can restart the watchdog timer at any time. The act of restarting is usually referred to as "kicking the dog". In this way, a program can be written which never lets the counter reach the terminal value.

4.1.1.2 Arduino Mega ADK

In case of an Arduino board, when the watchdog timer timeouts, the sketch is restarted (new call to *setup()*). Furthermore, a watchdog timeout signal is sent and this signal can be captured by the sketch. In the Arduino sketch, this is implemented by the JeeLib library (13). JeeLib is a library written for experimenting with JeeLabs products, however, some parts of the library are written for Arduino boards and can be used with them. Specifically, the Ports class (14) is the one used in this implementation to put the Arduino Mega ADK into sleep mode. The sketch execution stops for the specified sleep time and afterwards execution continues from the call to sleep function.

12. Add a sample overview of the sketch here maybe?

4.1 Power Consumption

4.1.1.3 WiFly module

The RN-XV WiFly module can be put to sleep in two ways - sleep timer or sleep command. With the sleep timer, the shield will enter sleep mode after a specified time period has passed since all TCP active connections have closed. With the sleep command the module will enter sleep mode immediately, unless an active TCP connection exists. (15).

This means that in order to put the WiFly module to sleep, all active connections must be closed. For the XMPP session, this means closing the active stream and the underlying TCP connection. Once this has been done, the module can successfully enter sleep mode.

The module can be waken up by either sending characters of the UART or by using the wake timer. In our implementation the activity on the UART wakes the module up when the sketch execution continues after the Mega ADK wakes up from sleep mode and establishes a new TCP connection in the start of *loop()* method call. This effectively means going through the first 5 stages of XMPP session lifecycle on every wake up. Since XMPP session initialization is quite verbose and the communication over Wi-Fi is unstable, the possibility of receiving scrambled or incomplete messages creates a problem.

4.1.2 Communication

Because all TCP connections and therefore the XMPP session have to be closed after every data transmit, XMPP session lifecycle steps 1 - 5 shown in Figure 3.1 are executed multiple times. When testing the XMPP implementation with Arduino Mega ADK and WiFly module sleep modes enabled, a troubling fact was discovered - the XMPP session negotiation fails at least once for every 30 minute test. The reason for these connection failures are scrambled authentication or stream opening stanzas.

From the tests, it could be seen that the average XMPP session start up time was 15 seconds. **13. show a graph here i guess? or whatever** Data needs to be transmitted every 10 seconds, which means that the module can never be put to sleep as it will not be able to go through the sleep and wake up cycle during the available time period. Moreover, data transmission took on average 1 second,

4.2 Server-side Client

meaning that 15 seconds spent on wake up would result in a second of actual work, which is not efficient.

During testing the sleep and wake up cycles with XMPP, it was found that XMPP session negotiation will hang approximately once in 15 minutes. As the maximum sleep time in our implementation is 65 seconds, this means that there are at least 13 separate session negotiations. Of course, this is the problem of the XMPP library in use and its lack of error handling. Problems with the library could be addressed with a better implementation, however, there was another factor discovered during the tests.

As a result, XMPP as a means of communication was not viable when trying to minimize power consumption. As an alternative, [14. look how to reference these](#)

web sockets, raw sockets and HTTP was considered. Web sockets were left out because opening a web socket connection is opened with a HTTP request, making using that one request to actually send the data more efficient. Therefore, HTTP was preferred to web sockets. Raw socket implementation in Arduino would add needless complexity to the sketch and was therefore not implemented.

The main advantages of HTTP are connection initialization speed and simplicity. The average time to wake up from sleep mode, send an HTTP request and receive response from the server, was measured to be around 6 seconds. In the previous scenario of 10 second transmission interval, it would mean 4 seconds could be spent in sleep mode, 4 seconds to wake up and transmit the data. This implementation would be more energy efficient, which can be seen in Figure 4.1. The data was gathered while running the prototype for 30 minutes with a data sending interval of 30 seconds. The sleep times were adjusted for both configurations based on their wake up times. As seen in the figure, the HTTP configuration consumes approximately 20% less power than the one using XMPP.

[15. Add a sample request here](#)

4.2 Server-side Client

With the move from XMPP to HTTP, the server-side client's implementation changed. Instead of using XMPP Smack library, a web server was needed. Jetty

4.2 Server-side Client




Figure 4.1: Power Consumption at 30 Second Interval

was selected because of its simplicity and possibility to embed it into the application. A web server embedded in an application is useful when prototyping, because it saves time on configuration and deployment.

Furthermore, to take full advantage of the newly developed sleep mode functionality, the client was further developed to predict sensor values for some time. Two modules were added to the server for this - simple linear regression and fuzzy control engine. In addition, the data model was modified to suit the new modules. The new data model can be seen in Figure 4.2. *Data* table now has *measured* field, which indicates if the value has been measured or predicted. *SensorTypes* table has two new fields - *regression_error* and *measure_error*. *regression_error* is the acceptable regression model error for future predictions. *measure_error* indicates the acceptable variance for measured values.

16. Do i really need this here? The final web server consists of 4 modules:

1. Request handler
2. Linear regression model
3. Fuzzy control engine

4.2 Server-side Client

4. Data storage




Figure 4.2: Updated Server-side Data Model

4.2.1 Simple Linear Regression

A simple linear regression model was selected to predict future data, which uses the least squares method to calculate the future values. 17. add a mention of the paper the idea was taken from? The model has a single explanatory variable - Unix time. This variable is used to predict the future values of sensor readings based on previous measurements.

The model sample is taken from previous measurements during the last 4 minutes. The 4 minute interval is selected to balance out errors caused by extreme values. This is necessary because a single value with big enough deviation can cause the regression model to become inaccurate.

To measure the accuracy of the model, an confidence level of 90% was introduced. If the sample for last 4 minutes provides an accurate enough regression model, then the predictions can be used. Otherwise, fresh data should be queried and added to the model until the error threshold is satisfied. The equation to calculate regression model's confidence level α :

$$\alpha = 2\Phi\left(\frac{\epsilon}{STDEV[e]}\right) - 1.$$

Here Φ is the CDF (Cumulative Distribution Function) of residuals, ϵ is the allowed error (in the database model named as *regression_error* and *STDEV[e]* is the standard deviation of residuals (also known as standard error of estimate)).

18. look this part over, should be ok

Regression models are created separately for each sensor. Each model has a sample of previous measurements and the allowed error ϵ defined in the database. The confidence level α is calculated for each model and has to be over 90% which is the selected confidence threshold. If all regression models satisfy the confidence level, then the fuzzy control system is initialized with data from previous measurements and calculated confidence levels.

4.2.2 Fuzzy Logic Engine

The next step in predicting future sensor values is to calculate the time interval for the next measurement request. To calculate the time, a fuzzy control system was introduced to provide flexible decisions based on multiple input variables. The input variables are:

1. *temp*
2. *temp predictability*
3. *light*
4. *light predictability*
5. *hall*
6. *hall predictability*

Here each sensor has a pair of input variables - measured sensor value and sensor value predictability (regression model confidence level α). All these variables are mapped into fuzzy sets, a procedure call fuzzification. The predictability variables are mapped into sets which all have the same membership functions shown in figure Figure 4.3.

The fuzzy control system is initialized with a previous measurement for each sensor. This effectively sets the bounds of each sensor's fuzzy sets as seen in Figure 4.4. The *PREDICTABLE* fuzzy set's DOM (degree of membership) reaches maximum at a confidence level of 0.95 or 95% because for our prototype anything above that level is highly predictable.

4.2 Server-side Client

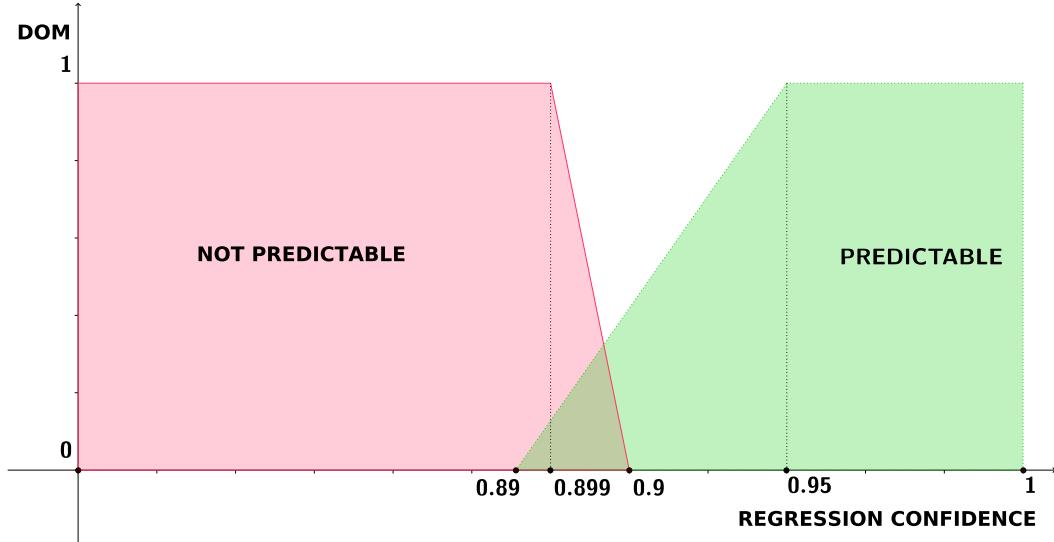


Figure 4.3: Regression Confidence Fuzzy Sets

The measured sensor values are fuzzified into sets, which each have different membership functions. The membership functions are calculated on the last measurement received as seen in Figure 4.4. Here X in set labels notes the sensor type currently used as the sets are the same relative to each sensor's previous measurement. The center point O is the previously measured value. $\|X_2X_3\|$ is the predefined sensor measurement error (*sensor_error* field in the data model). This means that if the new measurement is between the points X_2 and X_3 , there is no change in the measurement for the fuzzy system. The lengths $\|X_0X_2\|$, $\|X_1O\|$, $\|OX_4\|$ and $\|X_3X_5\|$ are defined by the same variable *slope_width*. Currently, slope width is the same as the selected *regression_error* values which are 5 times the value of *measure_error*. The selected error values can be seen in Table 4.1.

	Temp	Hall	LDR
<i>measure_error</i>	0.1	0.5	10
<i>regression_error</i>	0.5	2.5	50

Table 4.1: Selected Error Values

4.2 Server-side Client

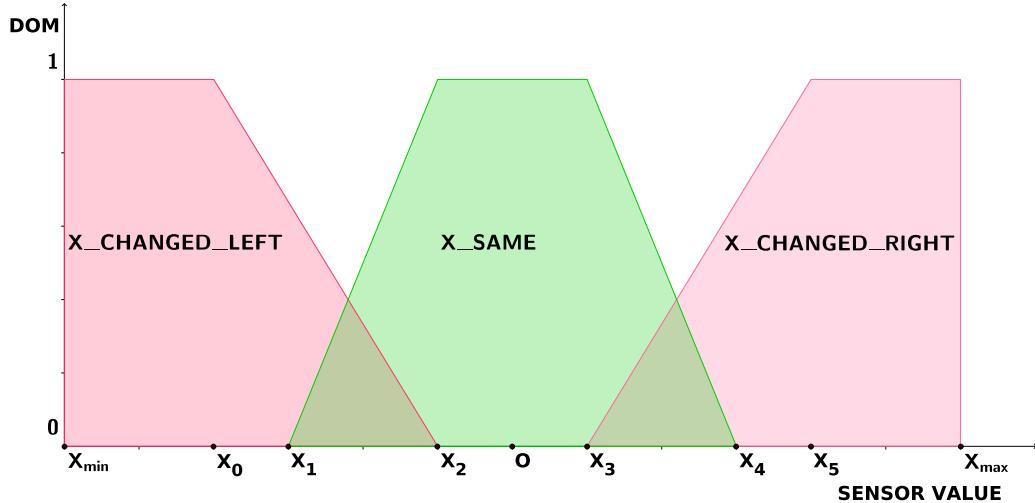


Figure 4.4: Sensor Value Fuzzy Sets

The output variable is defuzzified from the idle time sets which have membership functions described in Figure 4.5. The sets *request* and *predict* are symmetrical to enable easy defuzzification. The maximum idle time is 65 seconds, which is the maximum length of one sleep cycle for the sensor module.

When all input variables have been fuzzified and degrees of membership calculated, then the next step is to fire all predefined rules. For each sensor, there is a collection of 3 rules, making a total of 9 rules (here $x \in \{hall, temp, light\}$):

1. IF x IS *x_changed_left* OR *x_changed_right* AND *x_predictability* IS *not_predictable* OR *predictable* THEN *request*
2. IF x IS *x_same* AND *x_predictability* IS *not_predictable* THEN *request*
3. IF x IS *x_same* AND *x_predictability* IS *predictable* THEN *predict*

The previously calculated degrees of membership are fed into these rules. For *AND* relationship, the minimum value is selected, for *OR* relationship, the maximum value is selected. For each rule, the output variable received a DOM equal to the DOM of the premise. When all rules have been evaluated, the




Figure 4.5: Idle Time Fuzzy Sets

results are defuzzified into a single crisp value, which is the idle time decision. This process is called defuzzification.

During defuzzification the DOMs of result sets *request* and *predict* for each rule are combined into a single DOM for the specified set. For this, *request* output has an *OR* relationship (maximum value) and the *predict* set has an *AND* relationship (minimum value). These relationships mean that the most inaccurate sensors would be represented in the final result. For example if *hall* has a DOM of 1.0 to *predict*, but *temp* has a DOM of 0.5 to *predict*, then the outcome would be a DOM of 0.5, because we need the output for least accurate sensors. For *request* this logic is reversed as the maximum DOM to *request* represents the output for the least accurate sensor.

The next step is to get the crisp result value. For this process, the idle time values for *predict* and *request* set membership functions at the specified DOM are calculated. x is found the equation $f(x) = DOM$, where x is the idle time and f is the membership function. *request* and *predict* membership functions are symmetric with the axis being in the point $x = 0$ as seen in Figure 4.5. The final defuzzified idle time is equal to $\max\{x_{predict} - x_{request}, 10\}$. Meaning that if the DOM to *request* was higher than to *predict*, then the default idle time of 10

4.3 Results

seconds is returned. Otherwise the value at the center point of the two results is returned.

When the crisp result is received from the fuzzy control system and idle time is longer than 10 seconds, future data is predicted for the idle period. This means that using the regression models, future values with 10 second intervals are inserted into the database for each sensor.

Finally, the idle time with a minimum value of 10 seconds is written as a JSON format string to the response body of the HTTP request. A sample response would look like:

19. add a sample response here

```
HTTP/1.1 200 OK
Content-Type: application/json; charset=utf-8
Content-Length: 12
Connection: close
Server: Jetty(7.2.0.v20101020)

{"idle": 34}
```

4.3 Results

20. Describe tests carried out

21. Make a test with a 9V battery

22. Graphs with improved power consumption figures on it, description of what changed and how/why **23. Graph with 10s idle time as the worst-case scenario here**

5

Conclusions

5.1 Conclusions

5.2 Summary of Contributions

6

Related Work

A solution for energy-efficient data collection for clustering-based wireless sensor networks has been suggested (16). This solution has wireless sensor clusters with a head node, which actively decides if new sensor data should be requested from the nodes or if they can be predicted. The communication with the central data collection point goes through cluster heads only, which reduces transfer overhead. This is similar to the current thesis in a sense that data is predicted or requested and then forwarded, with the exception that this prototype does not have clusters and extra complexity that comes with clustering. The general implementation idea has been taken from this solution and prediction algorithm details, because in both cases simple linear regression is used to predict sensor data.

24. improve

this?

25. check linear regression reference from that thesis

The previous prototype (1) is the main basis of this work as it was the starting point. The previous implementation used XMPP as the communication protocol and a fixed interval data collection. XMPP has several advantages over simple HTTP requests, mainly because the protocol supports message queues and security. However, actually useful security options **26. which was it?** (SSL/TLS) are unavailable due to their overhead and implementation complexity for Arduino boards. On the other hand, XMPP protocol has communication overhead which reduces its usability when fast connection establishment and data transmission is required.

A similar system of Arduino sensor module and a central data collection server was discussed in three blog posts (17, 18, 19). The posts focused on power consumption of Arduino boards, because similarly a 9V battery could not power the device long enough. The main focus point the the posts was hardware modifications to the boards by either using a more power-efficient Arduino board, [27.](#) [maybe reference arduino mini here?](#) replacing certain components of the board or making their own hardware using a breadboard. On the software side, a similar sleep and wake cycle solution was developed. The sleep and wake cycles were implemented using the internal watchdog timer similar to what is used in the JeeLib library's Port class (14).

7

Future Research Directions

Briefly indicate how your current research can be extended, some improvements, etc.

28. Different arduino/hardware configuration 29. Better prediction

8

Sisukokkuvõte

Eesti abstract...

Bibliography

- [1] Kaarel Hanson, Context Sensor Data on Demand for Mobile Users Supported by XMPP, Tartu, Estonia. v, 1, 13, 15, 16, 17, 30
- [2] Arduino - HomePage, <http://www.arduino.cc/>.
URL <http://www.arduino.cc/> 3
- [3] Arduino - ArduinoBoardADK, <http://arduino.cc/en/Main/ArduinoBoardADK>.
URL <http://arduino.cc/en/Main/ArduinoBoardADK> 4
- [4] Arduino - ArduinoWirelessShield, <http://arduino.cc/en/Main/ArduinoWirelessShield>.
URL <http://arduino.cc/en/Main/ArduinoWirelessShield> 5
- [5] RN-XV WiFly module - wire antenna - SparkFun electronics, <https://www.sparkfun.com/products/10822>.
URL <https://www.sparkfun.com/products/10822> 6
- [6] Tinkerkit introduction, <http://www.tinkerkit.com/introduction/>.
URL <http://www.tinkerkit.com/introduction/> 7
- [7] Thermistor, <http://www.tinkerkit.com/thermistor/>.
URL <http://www.tinkerkit.com/thermistor/> 7
- [8] Hall, <http://www.tinkerkit.com/hall/>.
URL <http://www.tinkerkit.com/hall/> 7
- [9] Light sensor, <http://www.tinkerkit.com/ldr/>.
URL <http://www.tinkerkit.com/ldr/> 7
- [10] Lotfi A. Zadeh, Fuzzy Sets. *Information and Control*, 1965. 9
- [11] Marcelo Godoy Simoes, Introduction to Fuzzy Control, 1610 Illinois Street Golden, Colorado 80401-1887 USA. 10
- [12] Jim Lamberson, Single and multistage watchdog timers (2012). 19
- [13] JeeLib: introduction, <http://jeelabs.net/pub/docs/jeelib/>.
URL <http://jeelabs.net/pub/docs/jeelib/> 19
- [14] JeeLib: port class reference, <http://jeelabs.net/pub/docs/jeelib/classPort.html>.
URL <http://jeelabs.net/pub/docs/jeelib/classPort.html> 19, 31
- [15] WiFly user guide - WiFly-RN-UM.pdf, <http://dlnmh9ip6v2uc.cloudfront.net/datasheets/Wireless/WiFi/WiFly-RN-UM.pdf>.
URL <http://dlnmh9ip6v2uc.cloudfront.net/datasheets/Wireless/WiFi/WiFly-RN-UM.pdf> 20
- [16] H. Jiang, S. Jin, C. Wang, Prediction or not? an energy-efficient framework for clustering-based data collection in wireless sensor networks, *Parallel and Distributed Systems, IEEE Transactions on* 22 (6) (2011) 1064–1071.
doi:10.1109/TPDS.2010.174. 30
- [17] An idiot and an arduino: Ruminations on power « ~jmoskie, <http://www.josephmoskie.com/?p=292>.
URL <http://www.josephmoskie.com/?p=292> 31
- [18] How to run an arduino (clone) on a 9V battery for weeks or even months, <http://hwstartup.wordpress.com/2013/03/11/how-to-run-an-arduino-on-a-9v-battery-for-weeks-or-months/>.
URL <http://hwstartup.wordpress.com/2013/03/11/how-to-run-an-arduino-on-a-9v-battery-for-weeks-or-months/> 31
- [19] Low-power arduino using the watchdog timer | fiz-ix, <http://www.fiz-ix.com/2012/11/low-power-arduino-using-the-watchdog-timer/>.
URL <http://www.fiz-ix.com/2012/11/low-power-arduino-using-the-watchdog-timer/> 31

ToDo

	P.
1. Write this part	1
2. mention peaktech power supply here?	2
3. add more chapters here when they're ready	2
4. add an image of the IDE here...	3
5. format the text a bit, remove describe what is EEPROM, add an image	5
6. add citations here...	7
7. Add simple linear regression here, too?	12
8. Add a small introduction of who wrote the previous thesis etc	13
9. how should i reference all the data collected?	17
10. reference it somehow	18
11. introduction	18
12. Add a sample overview of the sketch here maybe?	19
13. show a graph here i guess? or whatever	20
14. look how to reference these	21
15. Add a sample request here	21
16. Do i really need this here?	22
17. add a mention of the paper the idea was taken from?	23
18. look this part over, should be ok	23
19. add a sample response here	28
20. Describe tests carried out	28
21. Make a test with a 9V battery	28
22. Graphs with improved power consumption figures on it, description of what changed and how/why	28
23. Graph with 10s idle time as the worst-case scenario here	28

BIBLIOGRAPHY

24. improve this?	30
25. check linear regression reference from that thesis	30
26. which was it?	30
27. maybe reference arduino mini here?	31
28. Different arduino/hardware configuration	32
29. Better prediction	32